

BASIL C K

AI Engineer | Machine Learning · LLM Systems · Computer Vision · RAG Pipelines

basilck618@gmail.com | +91 8111974510 | basilbaasi.github.io | LinkedIn | GitHub | Kochi, Kerala, India

PROFILE

AI Engineer Intern building production AI systems spanning Computer Vision, LLM integration, RAG pipelines, and AI workflow orchestration. Currently developing an AI model evaluation and orchestration platform for testing multi-provider LLMs while delivering Computer Vision and Generative AI solutions at **Difinity Digital**. Published research at **IEEE ICNPPCV 2026** on hybrid deep learning frameworks. Passionate about scalable AI infrastructure — from model experimentation to deployment-ready intelligent systems.

TECHNICAL SKILLS

AI & Machine Learning: Computer Vision, Deep Learning, CNNs, YOLO, LLM Integration, RAG, Prompt Engineering, Model Evaluation & Benchmarking, AI Workflow Orchestration, Conversational AI

LLM & Generative AI: Commercial LLM APIs (Gemini, OpenAI, NVIDIA), SSE Streaming, AI Memory Systems

Frameworks & Libraries: TensorFlow/Keras, PyTorch, OpenCV, FastAPI, Flask, SQLAlchemy, Pydantic, NumPy, Pandas

Infrastructure & Tools: Python, SQL, Docker, Docker Compose, Git, Linux, PostgreSQL, VS Code, Jupyter Notebook

Cloud & Platforms: Oracle Cloud Infrastructure (OCI AI Foundations), AWS Fundamentals, Google AI Essentials

EXPERIENCE

AI Engineer Intern — Difinity Digital

06/2026 – Present

Kochi, Kerala

- Developing AI-powered applications using Python, FastAPI, and REST APIs for real-world client deployments. Internship concludes **September 8, 2026**.
- Building Computer Vision solutions and Generative AI integrations including LLM-based workflow automation.
- Designing production-ready backend services and AI inference pipelines while collaborating with cross-functional engineering teams.

PROJECTS

OpenMind AI Platform — AI Model Evaluation & LLM Orchestration Platform **ONGOING**

07/2026 – Present

[Python](#) | [FastAPI](#) | [Gemini AI](#) | [LLM APIs](#) | [RAG](#) | [pgvector](#) | [SQLAlchemy](#) | [Docker](#) | [PostgreSQL](#)

- Building an AI model evaluation platform to test, compare, and benchmark LLMs across API-driven providers (Gemini, GPT-4o) with planned expansion to locally-hosted open-weight models (Llama 3, DeepSeek R1) via Ollama/vLLM.
- Designed a multi-session AI playground with configurable inference parameters (temperature, top_p, max_tokens, system prompts) and real-time SSE streaming for interactive model response evaluation.
- Engineered the foundation for a RAG pipeline featuring PDF/Markdown document ingestion and text chunking, and currently architecting the pgvector infrastructure for upcoming semantic knowledge retrieval.
- Architecting a graph-based multi-tier memory system (Conversation, Semantic, Long-Term) to enable persistent AI context across sessions — a core building block for agentic AI applications.
- Designed the architecture and simulation model for a multi-step AI workflow orchestration engine, paving the way for automated LLM routing, conditional branching, and API integration.
- Engineered an OpenAI-compatible API gateway with Bearer token authentication, enabling standardized model inference routing across 40+ API endpoints with structured logging and monitoring.

Hybrid CNN–YOLO Framework for Car Damage Severity Classification and Localization

2025 – 2026

Python | TensorFlow/Keras | YOLO | OpenCV | FastAPI | Computer Vision

- Developed an end-to-end hybrid deep learning framework integrating YOLO-based object detection with CNN-based severity classification for automated vehicle damage assessment.
- Built the complete ML pipeline including dataset preparation, model training, hyperparameter tuning, evaluation metrics, and a FastAPI inference backend for real-time predictions.
- Presented the research at the 2026 IEEE International Conference on Natural Language Processing and Computer Vision (ICNPCV), Bengaluru, India.

Vocalix — AI Voice Assistant

2025

Python | Speech Recognition | LLM APIs | Automation Libraries

- Built a voice-controlled AI assistant combining speech recognition with LLM inference for natural language task automation and system-level command execution.
- Engineered a real-time speech-to-LLM-to-action pipeline with modular command processing and dynamic response generation.

PUBLICATIONS

A Hybrid CNN–YOLO Framework for Car Damage Severity Classification and Localization

2026 IEEE International Conference on Natural Language Processing and Computer Vision (ICNPCV), Bengaluru, India

- Published and presented peer-reviewed research on a hybrid deep learning framework combining YOLO object detection with CNN classification for automated vehicle damage assessment.

EDUCATION

B.Tech in Artificial Intelligence and Data Science

2022 – 2026

JCT College of Engineering and Technology, Anna University

(Graduated)

CGPA: 8.47 / 10

CERTIFICATIONS

- Oracle Cloud Infrastructure 2023 — AI Foundations Associate
- Google AI Essentials
- IBM Data Science — Selected Coursework (Python, Data Analysis, Machine Learning)

LANGUAGES

English · Malayalam · Tamil